

Diamonds

#Diamonds at certain prices

```
subset(diamonds, price<500)
```

```
## # A tibble: 1,729 x 10
```

```
##   carat      cut color clarity depth table price      x      y      z
##   <dbl>    <ord> <ord>   <ord> <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  0.23    Ideal   E     SI2   61.5   55   326   3.95   3.98   2.43
## 2  0.21  Premium   E     SI1   59.8   61   326   3.89   3.84   2.31
## 3  0.23     Good   E     VS1   56.9   65   327   4.05   4.07   2.31
## 4  0.29  Premium   I     VS2   62.4   58   334   4.20   4.23   2.63
## 5  0.31     Good   J     SI2   63.3   58   335   4.34   4.35   2.75
## 6  0.24 Very Good   J    VVS2   62.8   57   336   3.94   3.96   2.48
## 7  0.24 Very Good   I    VVS1   62.3   57   336   3.95   3.98   2.47
## 8  0.26 Very Good   H     SI1   61.9   55   337   4.07   4.11   2.53
## 9  0.22     Fair   E     VS2   65.1   61   337   3.87   3.78   2.49
## 10 0.23 Very Good   H     VS1   59.4   61   338   4.00   4.05   2.39
## # ... with 1,719 more rows
```

```
subset(diamonds, price<250)
```

```
## # A tibble: 0 x 10
```


```
## # ... with 10 variables: carat <dbl>, cut <ord>, color <ord>,
## #   clarity <ord>, depth <dbl>, table <dbl>, price <int>, x <dbl>,
## #   y <dbl>, z <dbl>
```

```
subset(diamonds, price>=15000)
```

```
## # A tibble: 1,656 x 10
```

```
##   carat      cut color clarity depth table price      x      y      z
##   <dbl>    <ord> <ord>   <ord> <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  1.60    Ideal   G     VS2   61.9   56 15000   7.53   7.47   4.64
## 2  1.54  Premium   E     VS2   62.3   58 15002   7.31   7.39   4.58
## 3  1.19    Ideal   F    VVS1   61.5   55 15005   6.82   6.84   4.20
## 4  2.10  Premium   I     SI1   61.5   57 15007   8.25   8.21   5.06
## 5  1.69    Ideal   D     SI1   60.8   57 15011   7.69   7.71   4.68
## 6  1.50 Very Good   G    VVS2   62.9   56 15013   7.22   7.32   4.57
## 7  1.73 Very Good   G     VS1   62.8   57 15014   7.57   7.72   4.80
## 8  2.02  Premium   G     SI2   63.0   59 15014   8.05   7.95   5.03
## 9  2.05 Very Good   F     SI2   61.9   56 15017   8.13   8.18   5.05
## 10 1.50 Very Good   F     VS1   61.6   58 15022   7.35   7.43   4.55
## # ... with 1,646 more rows
```


```
#Exploring the largest peak area in the price histogram  
ggplot(diamonds, aes(x=price)) +geom_histogram(color='black', binwidth=10)+  
  coord_cartesian(xlim = c(250,925))
```



#Creating histogram of diamond prices by cut

```
ggplot(data = diamonds)+  
  geom_histogram(aes(x = price)) +  
  facet_wrap(~ cut)
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



#Getting Median stats by cut

```
aggregate(. ~cut, data=diamonds, FUN=median)
```

```
##      cut carat color clarity depth table price      x      y      z
## 1 Fair   1.00    4      3  65.0    58 3282.0  6.175  6.10  3.97
## 2 Good   0.82    3      3  63.4    58 3050.5  5.980  5.99  3.70
## 3 Very Good 0.71    3      4  62.1    58 2648.0  5.740  5.77  3.56
## 4 Premium 0.86    4      4  61.4    59 3185.0  6.110  6.06  3.72
## 5 Ideal  0.54    4      4  61.8    56 1810.0  5.250  5.26  3.23
```

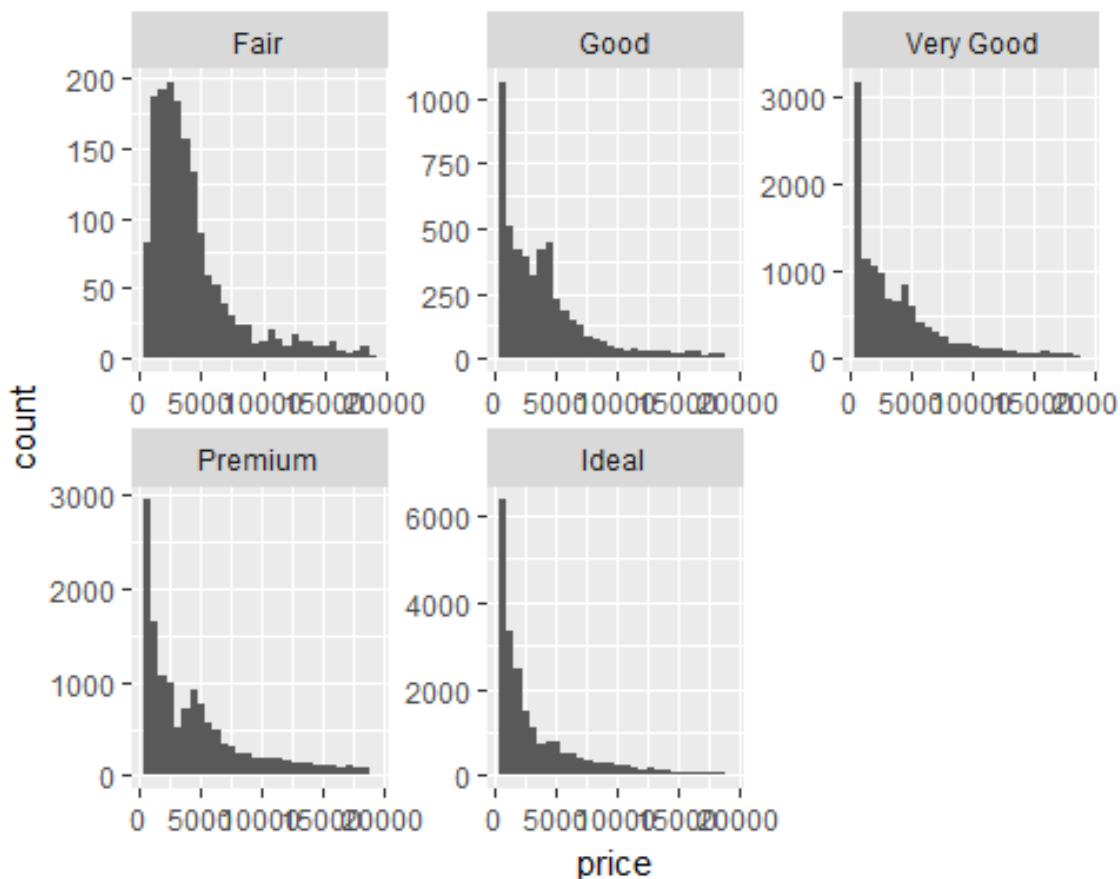
#Getting Highest price diamond by cut

```
aggregate(. ~cut, data=diamonds, FUN=max)
```

```
##      cut carat color clarity depth table price      x      y      z
## 1 Fair   5.01    7      8  79.0    95 18574 10.74 10.54  6.98
## 2 Good   3.01    7      8  67.0    66 18788  9.44  9.38  5.79
## 3 Very Good 4.00    7      8  64.9    66 18818 10.01  9.94 31.80
## 4 Premium 4.01    7      8  63.0    62 18823 10.14  9.90  8.06
## 5 Ideal  3.50    7      8  66.7    63 18806  9.65  9.60  6.03
```

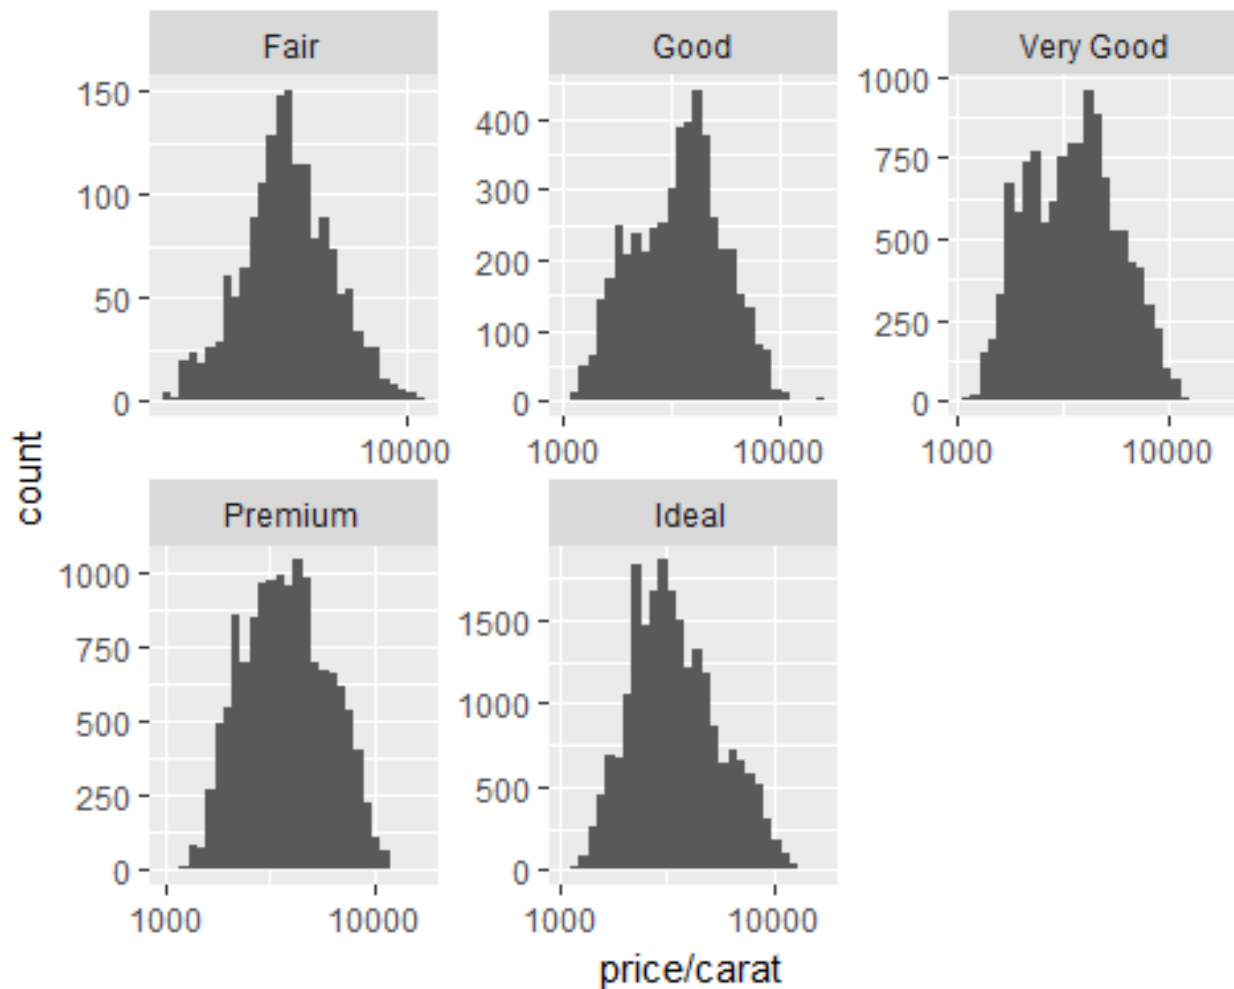
#The previous plot showed the distributions as different so freeing scales

```
ggplot(data = diamonds)+
  geom_histogram(aes(x = price)) +
  facet_wrap(~ cut, scales='free')
```

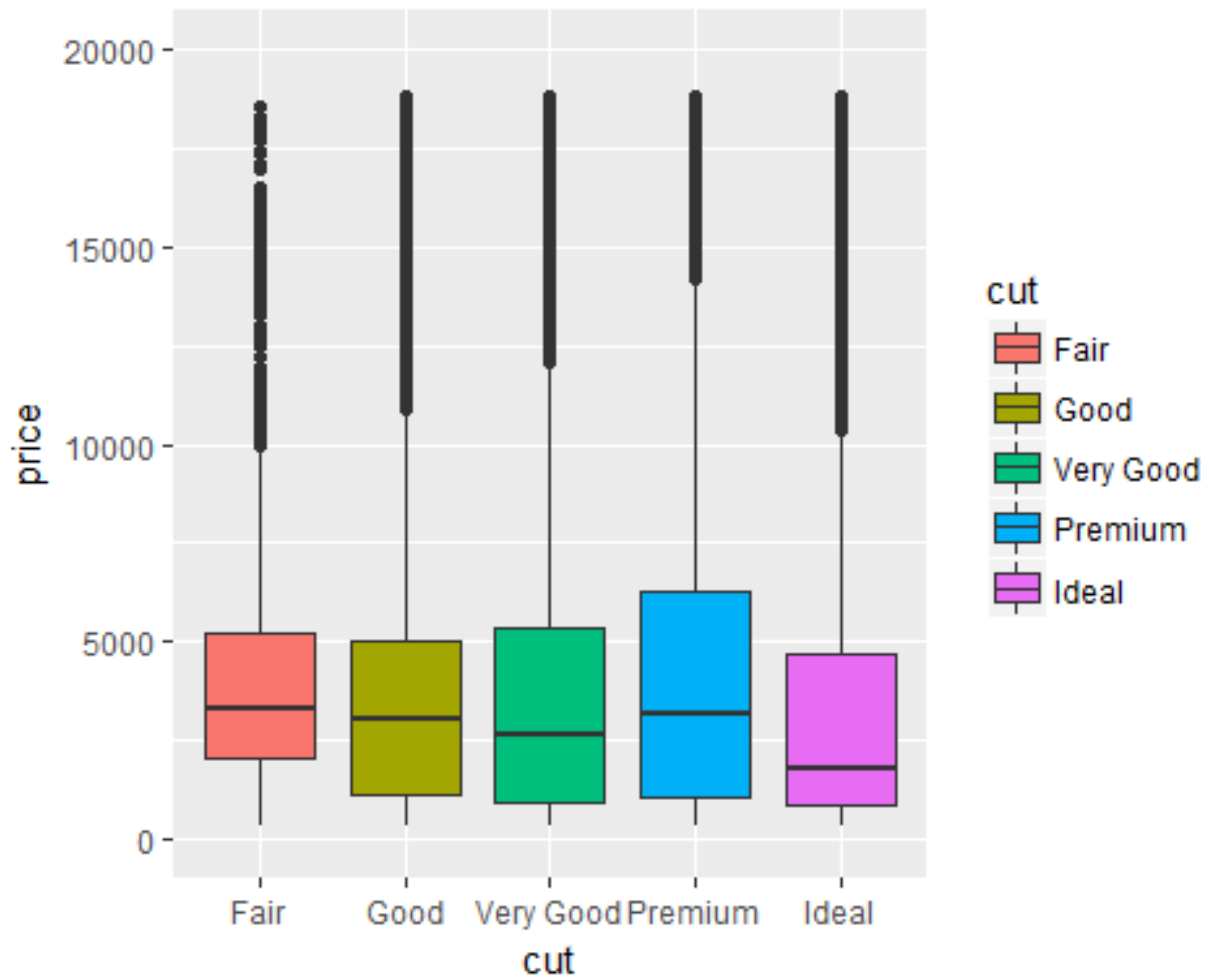


#Looking at the price per carat with log10

```
ggplot(data = diamonds)+  
  geom_histogram(aes(x = price/carat)) +  
  facet_wrap(~ cut, scales='free')+  
  scale_x_log10()
```



```
#Creating a box plot for price based on cut  
ggplot(diamonds, aes(x=cut, y=price, fill=cut))+  
  geom_boxplot()+  
  coord_cartesian(ylim = c(0,20000))
```



#Summarizing price data by color

```
by(diamonds$price, diamonds$color, summary)
```

```
## diamonds$color: D
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	357	911	1838	3170	4214	18693

```
## -----
```

```
## diamonds$color: E
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	326	882	1739	3077	4003	18731

```
## -----
```

```
## diamonds$color: F
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	342	982	2344	3725	4868	18791

```
## -----
```

```
## diamonds$color: G
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	354	931	2242	3999	6048	18818

```
## -----
```

```
## diamonds$color: H
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	337	984	3460	4487	5980	18803

```
## -----
```

```
## diamonds$color: I
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	334	1120	3730	5092	7202	18823

```
## -----
```

```
## diamonds$color: J
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	335	1860	4234	5324	7695	18710

#Getting individual IQR's based on certain color

```
IQR(subset(diamonds, color=='D')$price)
```

```
## [1] 3302.5
```

```
IQR(subset(diamonds, color=='J')$price)
```

```
## [1] 5834.5
```

#Getting price IQR by color

```
by(diamonds$price,diamonds$color,IQR)
```

```
## diamonds$color: D
```

```
## [1] 3302.5
```

```
## -----
```

```
## diamonds$color: E
```

```
## [1] 3121
```

```
## -----
```

```
## diamonds$color: F
```

```
## [1] 3886.25
```

```
## -----
```

```
## diamonds$color: G
```

```
## [1] 5117
```

```
## -----
```

```
## diamonds$color: H
```

```
## [1] 4996.25
```

```
## -----
```

```
## diamonds$color: I
```

```
## [1] 6081.25
```


```
## -----
```

```
## diamonds$color: J
```

```
## [1] 5834.5
```

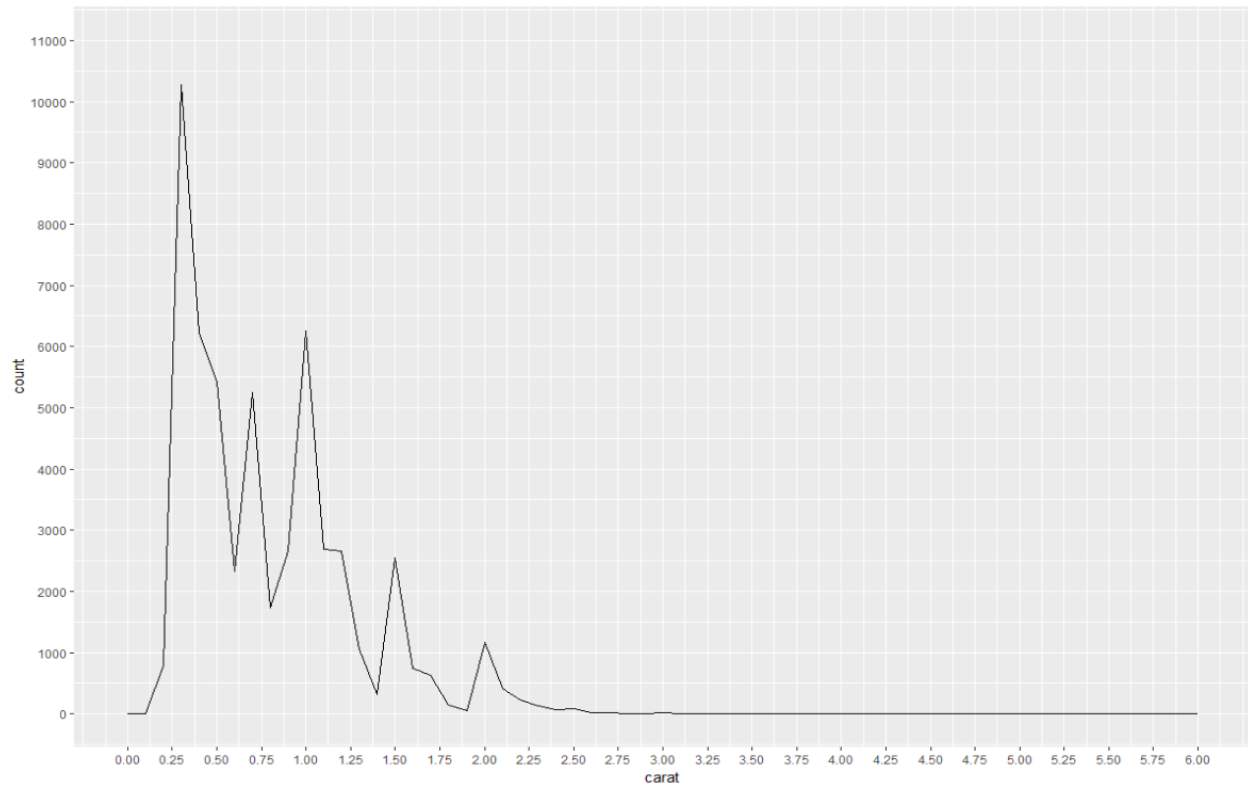


```
#Looking at price per carat of diamonds across colors  
ggplot(diamonds, aes(x=color, y=price/carat, fill=color))+  
  geom_boxplot()+  
  coord_cartesian(ylim = c(1000,7000))
```




#Investigating carats using a frequency polygon

```
ggplot(diamonds, aes(x=carat))+  
  geom_freqpoly(binwidth=0.1)+  
  scale_x_continuous(limits=c(0,6), breaks=seq(0,6,0.25))+  
  scale_y_continuous(limits=c(0,11000), breaks=seq(0,11000,1000))
```



```
ggplot(diamonds, aes(x=x, y=price))+geom_point(color="blue", alpha=1/20)
```



```
#Calculating correlations
```

```
with(diamonds, cor.test(x, price))
```

```
## Pearson's product-moment correlation
```

```
##
```

```
## data: x and price
```

```
## t = 440.16, df = 53938, p-value < 2.2e-16
```

```
## alternative hypothesis: true correlation is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 0.8825835 0.8862594
```

```
## sample estimates:
```

```
## cor
```

```
## 0.8844352
```

```
with(diamonds, cor.test(y, price))
```

```
## Pearson's product-moment correlation
```

```
##
```

```
## data: y and price
```

```
## t = 401.14, df = 53938, p-value < 2.2e-16
```

```
## alternative hypothesis: true correlation is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 0.8632867 0.8675241
```

```
## sample estimates:
```

```
## cor
```

```
## 0.8654209
```


```

with(diamonds, cor.test(z, price))

##
## Pearson's product-moment correlation
##
## data: z and price
## t = 393.6, df = 53938, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.8590541 0.8634131
## sample estimates:
## cor
## 0.8612494

ggplot(diamonds, aes(x=depth, y=price)) +geom_point(color='blue',
alpha=1/100)+
  scale_x_continuous(breaks = seq(0,80,2))

```




```

with(diamonds, cor.test(x=depth, y=price))


##
## Pearson's product-moment correlation
##
## data: depth and price
## t = -2.473, df = 53938, p-value = 0.0134
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.019084756 -0.002208537
## sample estimates:
## cor
## -0.0106474

ggplot(diamonds, aes(x=carat, y=price))+geom_point(color='blue', alpha=1/20)+
  scale_x_continuous(limits = c(0, quantile(diamonds$carat, 0.99)))+
  scale_y_continuous(limits = c(0, quantile(diamonds$price, 0.99)))+
  stat_smooth(method = 'lm')

```




```
ggplot(diamonds, aes(x= x*y*z, y=price)) +geom_point(alpha = 1/20,  
color='blue')
```



```
#Calculating volume with variables x, y,z  
#Getting correlation but not including volume of 0 or 800+  
diamonds$volume <- diamonds$x * diamonds$y * diamonds$z  
with(subset(diamonds,diamonds$volume>0 & diamonds$volume<800),  
cor.test(volume, price))
```

```
##  
## Pearson's product-moment correlation  
##  
## data: volume and price  
## t = 559.19, df = 53915, p-value < 2.2e-16  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## 0.9222944 0.9247772  
## sample estimates:  
## cor  
## 0.9235455
```



```
ggplot(subset(diamonds,diamonds$volume>0 & diamonds$volume<800), aes(x=
x*y*z, y=price)) +
  geom_point(alpha = 1/50, color='blue')+
  geom_smooth(method = 'lm', color='red')
```



```
diamondsByClarity <- diamonds %>%
  group_by(clarity) %>%
  summarise(mean_price = mean(price),
            median_price = median(price),
            min_price = min(price),
            max_price = max(price),
            n = n()) %>%
  arrange(clarity)
```

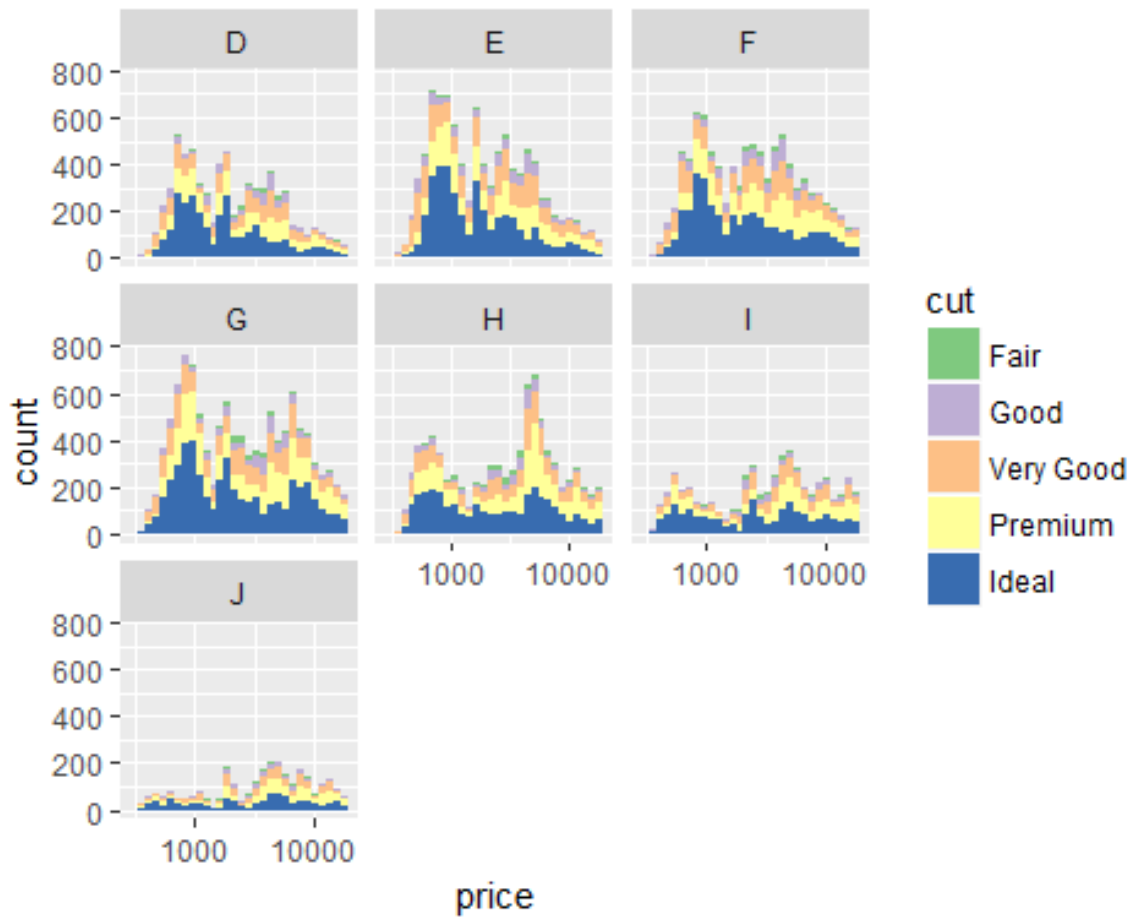
```
library(gridExtra)

b1 <- ggplot(diamondsByClarity, aes(x=clarity, y=mean_price))+
  geom_bar(stat='identity', color='blue')
b2 <- ggplot(diamondsByClarity, aes(x=clarity, y=median_price))+
  geom_bar(stat='identity', color='blue')
grid.arrange(b1,b2)
```




#Plotting the histograms in each color with a facet wrap showing the different cuts

```
ggplot(diamonds, aes(x=price, fill=cut))+ geom_histogram()+  
facet_wrap(~color)+  
  scale_fill_brewer(type = 'qual')+  
  scale_x_log10()
```




```
#Creating a scatter plot of the table value vs price
#Table reflects the width of the top diamond
ggplot(diamonds, aes(x=table, y=price)) +
  geom_point(aes(color=cut), alpha=1/2) +
  scale_color_brewer(type='qual') +
  coord_cartesian(xlim=c(50,80))
```



```
#Adding a volume variable with transform
diamonds <- transform(diamonds, volume = x*y*z)

#Plotting price (in a log10 scale) in terms of volume colored by clarity
ggplot(subset(diamonds, volume < quantile(volume, 0.99)), aes(x=volume,
y=price))+
  geom_point(aes(color = clarity))+
  scale_color_brewer(type = 'div')+
  scale_y_log10()
```



```
ggplot(diamonds, aes(x=cut, y=price)) +geom_jitter(aes(color=color))+
  scale_color_brewer(type = 'div')+
  facet_wrap(~ clarity)
```




```

cuberoot_trans = function() trans_new('cuberoot', transform = function(x)
x^(1/3),
                                     inverse = function(x) x^3)


#Looking at the carat with price on a log10 scale
library(scales)
ggplot(aes(x = carat, y = price, color=clarity), data = diamonds) +
  geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div',
    guide = guide_legend(title = 'Clarity', reverse = T,
      override.aes = list(alpha = 1, size = 2))) +
  scale_x_continuous(trans = cuberoot_trans(), limits = c(0.2, 3),
    breaks = c(0.2, 0.5, 1, 2, 3)) +
  scale_y_continuous(trans = log10_trans(), limits = c(350, 15000),
    breaks = c(350, 1000, 5000, 10000, 15000)) +
  ggtitle('Price (log10) by Cube-Root of Carat and Clarity')

```

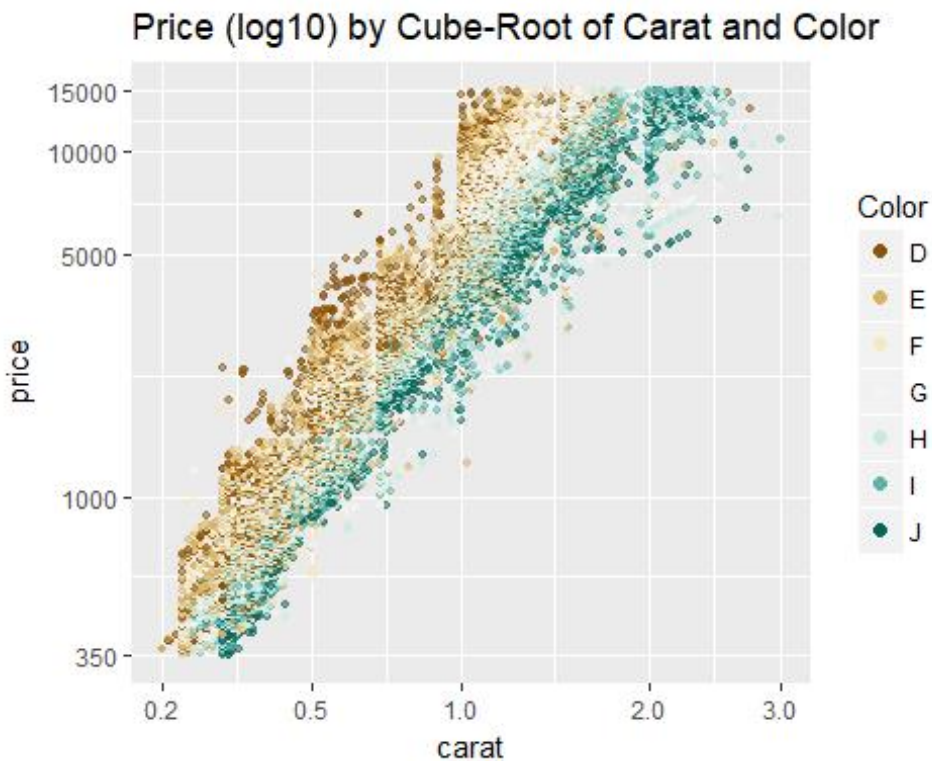


```
#Also looking at it colored by cut
ggplot(aes(x = carat, y = price, color = cut), data = diamonds) +
  geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div',
                    guide = guide_legend(title = 'Cut', reverse = T,
                                         override.aes = list(alpha = 1, size
= 2))) +
  scale_x_continuous(trans = cuberoot_trans(), limits = c(0.2, 3),
                    breaks = c(0.2, 0.5, 1, 2, 3)) +
  scale_y_continuous(trans = log10_trans(), limits = c(350, 15000),
                    breaks = c(350, 1000, 5000, 10000, 15000)) +
  ggtitle('Price (log10) by Cube-Root of Carat and Clarity')

## Warning: Removed 1693 rows containing missing values (geom_point).
```



```
ggplot(aes(x = carat, y = price, color = color), data = diamonds) +
  geom_point(alpha = 0.5, size = 1, position = 'jitter') +
  scale_color_brewer(type = 'div',
                    guide = guide_legend(title = 'Color', reverse = FALSE,
                                         override.aes = list(alpha = 1, size
= 2))) +
  scale_x_continuous(trans = cuberoot_trans(), limits = c(0.2, 3),
                    breaks = c(0.2, 0.5, 1, 2, 3)) +
  scale_y_continuous(trans = log10_trans(), limits = c(350, 15000),
                    breaks = c(350, 1000, 5000, 10000, 15000)) +
  ggtitle('Price (log10) by Cube-Root of Carat and Color')
```



#'I' makes it as is before using it in the regression, rather than as part of the formula
`suppressWarnings(library(memisc))`

```
m1 <- lm(I(log(price)) ~ I(carat^(1/3)), data= diamonds)
m2 <- update(m1, ~ . + carat)
m3 <- update(m2, ~ . + cut)
m4 <- update(m3, ~ . + color)
m5 <- update(m4, ~ . + clarity)
mtable(m1, m2, m3, m4, m5)
## Calls:
## m1: lm(formula = I(log(price)) ~ I(carat^(1/3)), data = diamonds)
## m2: lm(formula = I(log(price)) ~ I(carat^(1/3)) + carat, data = diamonds)
## m3: lm(formula = I(log(price)) ~ I(carat^(1/3)) + carat + cut, data = diamonds)
## m4: lm(formula = I(log(price)) ~ I(carat^(1/3)) + carat + cut + color,
##      data = diamonds)
## m5: lm(formula = I(log(price)) ~ I(carat^(1/3)) + carat + cut + color +
##      clarity, data = diamonds)
##
```

```

##
=====
##          m1          m2          m3          m4          m5
## -----
-
## (Intercept)      2.821***      1.039***      0.874***      0.932***      0.415***
##                (0.006)      (0.019)      (0.019)      (0.017)      (0.010)
## I(carat^(1/3))    5.558***      8.568***      8.703***      8.438***      9.144***
##                (0.007)      (0.032)      (0.031)      (0.028)      (0.016)
## carat            -1.137***      -1.163***      -0.992***      -1.093***
##                (0.012)      (0.011)      (0.010)      (0.006)
## cut: .L           0.224***      0.224***      0.120***
##                (0.004)      (0.004)      (0.002)
## cut: .Q          -0.062***      -0.062***      -0.031***
##                (0.004)      (0.003)      (0.002)
## cut: .C           0.051***      0.052***      0.014***
##                (0.003)      (0.003)      (0.002)
## cut: ^4           0.018***      0.018***      -0.002
##                (0.003)      (0.002)      (0.001)
## color: .L        -0.373***      -0.441***
##                (0.003)      (0.002)
## color: .Q        -0.129***      -0.093***
##                (0.003)      (0.002)
## color: .C         0.001         -0.013***
##                (0.003)      (0.002)
## color: ^4         0.029***      0.012***
##                (0.003)      (0.002)
## color: ^5        -0.016***      -0.003*
##                (0.003)      (0.001)
## color: ^6        -0.023***      0.001
##                (0.002)      (0.001)
## clarity: .L      0.907***
##                (0.003)
## clarity: .Q     -0.240***
##                (0.003)
## clarity: .C      0.131***
##                (0.003)
## clarity: ^4     -0.063***
##                (0.002)
## clarity: ^5      0.026***
##                (0.002)
## clarity: ^6     -0.002
##                (0.002)
## clarity: ^7      0.032***
##                (0.001)
## -----
-
## R-squared         0.924         0.935         0.939         0.951         0.984
## adj. R-squared    0.924         0.935         0.939         0.951         0.984
## sigma            0.280         0.259         0.250         0.224         0.129
## F                652012.063      387489.366      138654.523      87959.467      173791.084
## p                0.000         0.000         0.000         0.000         0.000
## Log-likelihood    -7962.499      -3631.319      -1837.416      4235.240      34091.272
## Deviance         4242.831      3613.360      3380.837      2699.212      892.214
## AIC              15930.999      7270.637      3690.832      -8442.481      -68140.544
## BIC              15957.685      7306.220      3761.997      -8317.942      -67953.736
## N                53940         53940         53940         53940         53940
##
=====

```